

AdvisorMatch

A Prototype System for
Streamlining Thesis Advisor
Discovery

Amit Kumar - 237007553
Archit Panwar - 533000200
Harshith Reddy - 537000715
Sanjit Sahu - 737002703

Motivation

High Stakes, Outdated Process

Critical career decision relying on obsolete, manual navigation of static directories.

Gap Exists

Semantic search exists for papers, not advisors; directories limited to rigid keyword matching.

The Efficiency Bottleneck

Fragmented workflow; hours wasted cross-referencing university profiles with external databases.

Need for Advisor-Centric Search

Current tools find papers, not mentors; need faculty ranking by aggregate expertise.

Current Process

Visit faculty directory

Check faculty interests

Search papers

Filter by relevance

Manually rank professors

Existing Systems

Academic Search Engines

e.g., Semantic Scholar

Limitation: Paper-Centric.

Great for documents, bad for people. No "advisor-view."

Requires manual aggregation to check active niches.

Generative AI

e.g., Gemini / ChatGPT

Limitation: Reliability & Cost.

Prone to hallucinations. High latency. Expensive. Lacks real-time verification.

Metric-Based Rankings

e.g., CSRankings

Limitation: Volume over Relevance.

Counts publications, ignores topic semantics. Answers "Who publishes most?" not "Who fits my research?"

Our Proposal: Semantic Advisor Search with Relevance Ranking

Intuitive Input

Accepts natural language queries describing research interests (e.g., "privacy in federated learning")—no rigid keywords required.

Semantic Analysis

Uses Vector Search (Sentence-BERT) to understand the meaning of the query and match it against faculty publication abstracts.

Advisor-Centric Rank

Aggregates paper relevance to score the professor, prioritizing recent work and consistent research activity.

Smart Results

Generates an ordered list of faculty members, instantly identifying the best candidates based on overall alignment.

Contextual Evidence

Highlights the specific publications that triggered the match, providing transparency and proof of the advisor's relevance.

Data collection, Annotation and Preprocessing

1. Scrape Faculty Name, Image, Interests, etc from Texas A&M CSE directory
2. Use OpenAlex API: search for faculty → get faculty ID →get papers.
3. Create a database of faculties and paper mappings with metadata (year of publication, citation count, first author, etc.)
4. Generate embeddings for paper abstract.
5. Build FAISS Index → Create search index.
6. Manually curate a small list of annotated results for testing. We used MRR and NDCG to analyze most relevant professor per query.

Data Statistics

Data Relevance

80.9%

of papers published
2020–2025.



Takeaway:

Matching is based on current
research trends, not outdated
history.

 **1,847**

Total Publications

 **85**

Faculty Coverage


TAMU CSE Professors

 **11,506**

Research Impact

Total Citations


Research Topics




Count | Percentage

Count | Percentage

Implementation



Technology Stack



Pre - Processing Step

Scrape Professors - Ingest
script scraping from
university websites

OpenAlex Author Search -
Finding and matching
professors with Author names
and with tamu affiliations

Database Storage - Storing
everything in SQLite

Generate Embeddings -
Creating semantic embeddings
with Sentence-BERT

Fetch Papers from OpenAlex
- Getting all papers and
metadata

Build FAISS Index - Building
the search index

Implementation - Initial Paper Ranking

FAISS Similarity Search

Index Type: FAISS IndexFlatIP (Inner Product)

Method: Cosine similarity on normalized vectors

Retrieval: Top 50 papers [TOP_K_PAPERS = 50]

Performance: Sub-100ms search time

Advantages:

Fast: Sub-100ms for 1,847 papers

Scalable: Efficient approximate nearest neighbor search

Semantic: Understands meaning, not just keywords

Search Process

Query Embedding (384-dim)

FAISS Index Search

Top 50 Papers + Similarity Scores (0.0 to 1.0)

Implementation - Re ranking

Paper Aggregation

- Group top 50 papers by professor
- Extract: similarity, year, author_position

Multi-Factor Scoring

Weighted Similarity

- Formula: $\text{similarity} \times \text{recency_weight}$
- First-author bonus: $\times 1.2$ (20% boost)
- Second-author bonus: $\times 1.1$ (10% boost)
- Average across top 10 papers

Recency Weight (Exponential Decay)

- Formula: $\exp(-0.1 \times \text{years_ago})$
- Recent papers weighted higher
- Example: 2024 paper = 0.905, 2020 paper = 0.607

Activity Bonus

- Formula: $\min(\text{recent_papers} \times 0.05, 0.20)$
- Rewards professors with recent publications (2022-2025)
- Capped at 0.20 (4+ recent papers)

Citation Impact

- Formula: $\log_{10}(1 + \text{citations}) / 3.0 \times 0.15$
- Log-normalized citation scores
- Weighted contribution: 15% of final score

Final Score =
 $\text{avg}(\text{weighted_scores}) + \text{activity_bonus} + \text{citation_impact}$

Implementation - Re ranking

$$\text{Score}_p = \text{AvgSim}_p + \text{Activity}_p + \text{Citation}_p$$

1. Average Weighted Similarity (AvgSim_p)

$$\text{AvgSim}_p = (1 / 10) \cdot \sum_{i=1}^{10} (\text{sim}_i \cdot \exp(-0.1 \cdot \text{yearsAgo}_i) \cdot \text{bonus}_i)$$

- i : one of professor p's top 10 papers
- sim_i : semantic similarity between the query and paper i
- yearsAgo_i : how many years ago paper i was published
- exp(-0.1 · yearsAgo_i) : recency weight (recent papers > old papers)
- bonus_i : 1.2 if first author, 1.1 if second author, 1.0 otherwise

3. Citation Impact (Citation_p)

$$\text{Citation}_p = [\log_{10}(1 + \text{citations}_p) / 3.0] \cdot 0.15$$

- citations_p : total citations for p's papers
- log₁₀(1 + citations_p) : compresses very large citation counts
- Division by 3.0 normalizes; × 0.15 => about 15% of final score

2. Activity Bonus (Activity_p)

$$\text{Activity}_p = \min(0.05 \cdot \text{recentPapers}_p, 0.20)$$

- recentPapers_p : number of papers by p from 2022–2025
- Adds 0.05 per recent paper, capped at 0.20 (4+ recent papers)

Variable Legend

- p : a professor
- i : one of p's papers considered in scoring
- Score_p : final advisor score used for ranking


Results and Model Comparison

For Queries- Use complex semantic queries that test conceptual understanding

Metric	BM25	Embedding	Improvement
MRR	0.3611	0.3750	+3.8%
NDCG@3	0.3977	0.4401	+10.7%
NDCG@4	0.3977	0.4401	+10.7%

A dark, grainy photograph of a stadium with many empty seats. The seats are arranged in rows, creating a pattern of light and shadow. The perspective is from the bottom left, looking up at the stands. The overall mood is somber and quiet.

Demo



Thank You